

# Insurance Claims Analysis

May 8, 2023

## 1 Insurance Claims Analysis - Group Project ACT SCI 657

The data is sourced from Kaggle, it can be accessed through the following link:

<https://www.kaggle.com/datasets/bunttyshah/auto-insurance-claims-data?datasetId=45152&sortBy=voteCount>

```
[ ]: # Import the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import \
    accuracy_score, confusion_matrix, classification_report
```

## 2 Data Reading

```
[ ]: # Read the xls file
df=pd.read_excel('insurance_claims.xlsx')
```

```
[ ]: # Preview the data
df.head()
```

```
[ ]:  months_as_customer  age  policy_number  policy_bind_date  policy_state  \
0          328      48        521585      2014-10-17          OH
1          228      42        342868      2006-06-27          IN
2          134      29        687698      2000-09-06          OH
3          256      41        227811      1990-05-25          IL
4          228      44        367455      2014-06-06          IL

    policy_cs1  policy_deductable  policy_annual_premium  umbrella_limit  \
0    250/500          1000          1406.91          0
1    250/500          2000          1197.22      5000000
```

2	100/300	2000	1413.14	5000000
3	250/500	2000	1415.74	6000000
4	500/1000	1000	1583.91	6000000

	insured_zip	...	police_report_available	total_claim_amount	injury_claim	\
0	466132	...	YES	71610	6510	
1	468176	...	?	5070	780	
2	430632	...	NO	34650	7700	
3	608117	...	NO	63400	6340	
4	610706	...	NO	6500	1300	

	property_claim	vehicle_claim	auto_make	auto_model	auto_year	\
0	13020	52080	Saab	92x	2004	
1	780	3510	Mercedes	E400	2007	
2	3850	23100	Dodge	RAM	2007	
3	6340	50720	Chevrolet	Tahoe	2014	
4	650	4550	Accura	RSX	2009	

	fraud_reported	_c39
0	Y	NaN
1	Y	NaN
2	N	NaN
3	Y	NaN
4	N	NaN

[5 rows x 40 columns]

```
[ ]: # List the columns
df.columns
```

```
[ ]: Index(['months_as_customer', 'age', 'policy_number', 'policy_bind_date',
'policy_state', 'policy_csl', 'policy_deductable',
'policy_annual_premium', 'umbrella_limit', 'insured_zip', 'insured_sex',
'insured_education_level', 'insured_occupation', 'insured_hobbies',
'insured_relationship', 'capital-gains', 'capital-loss',
'incident_date', 'incident_type', 'collision_type', 'incident_severity',
'authorities_contacted', 'incident_state', 'incident_city',
'incident_location', 'incident_hour_of_the_day',
'number_of_vehicles_involved', 'property_damage', 'bodily_injuries',
'witnesses', 'police_report_available', 'total_claim_amount',
'injury_claim', 'property_claim', 'vehicle_claim', 'auto_make',
'auto_model', 'auto_year', 'fraud_reported', '_c39'],
dtype='object')
```

### 3 Exploratory Data Analysis

```
[ ]: # Summary the numerical columns
df.describe()
```

```
[ ]:      months_as_customer      age  policy_number  policy_deductable \
count      1000.000000    1000.000000    1000.000000    1000.000000
mean        203.954000    38.948000   546238.648000    1136.000000
std         115.113174     9.140287   257063.005276     611.864673
min           0.000000    19.000000   100804.000000     500.000000
25%         115.750000    32.000000   335980.250000     500.000000
50%         199.500000    38.000000   533135.000000    1000.000000
75%         276.250000    44.000000   759099.750000    2000.000000
max         479.000000    64.000000   999435.000000    2000.000000

      policy_annual_premium  umbrella_limit  insured_zip  capital-gains \
count      1000.000000    1.000000e+03    1000.000000    1000.000000
mean        1256.406150    1.101000e+06   501214.488000    25126.100000
std          244.167395    2.297407e+06    71701.610941    27872.187708
min          433.330000   -1.000000e+06   430104.000000     0.000000
25%         1089.607500    0.000000e+00   448404.500000     0.000000
50%         1257.200000    0.000000e+00   466445.500000     0.000000
75%         1415.695000    0.000000e+00   603251.000000    51025.000000
max         2047.590000    1.000000e+07   620962.000000   100500.000000

      capital-loss  incident_hour_of_the_day  number_of_vehicles_involved \
count      1000.000000    1000.000000    1000.000000
mean     -26793.700000     11.644000     1.83900
std      28104.096686     6.951373     1.01888
min     -111100.000000     0.000000     1.00000
25%     -51500.000000     6.000000     1.00000
50%     -23250.000000    12.000000     1.00000
75%       0.000000    17.000000     3.00000
max       0.000000    23.000000     4.00000

      bodily_injuries  witnesses  total_claim_amount  injury_claim \
count      1000.000000    1000.000000    1000.000000    1000.000000
mean         0.992000    1.487000    52761.94000    7433.420000
std         0.820127    1.111335    26401.53319    4880.951853
min         0.000000    0.000000    100.00000     0.000000
25%         0.000000    1.000000    41812.50000    4295.000000
50%         1.000000    1.000000    58055.00000    6775.000000
75%         2.000000    2.000000    70592.50000   11305.000000
max         2.000000    3.000000   114920.00000   21450.000000

      property_claim  vehicle_claim  auto_year  _c39
count      1000.000000    1000.000000    1000.000000    0.0
```

mean	7399.570000	37928.950000	2005.103000	NaN
std	4824.726179	18886.252893	6.015861	NaN
min	0.000000	70.000000	1995.000000	NaN
25%	4445.000000	30292.500000	2000.000000	NaN
50%	6750.000000	42100.000000	2005.000000	NaN
75%	10885.000000	50822.500000	2010.000000	NaN
max	23670.000000	79560.000000	2015.000000	NaN

```
[ ]: # Info the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   months_as_customer                    1000 non-null   int64
1   age                                   1000 non-null   int64
2   policy_number                         1000 non-null   int64
3   policy_bind_date                     1000 non-null   datetime64[ns]
4   policy_state                          1000 non-null   object
5   policy_csl                           1000 non-null   object
6   policy_deductable                     1000 non-null   int64
7   policy_annual_premium                 1000 non-null   float64
8   umbrella_limit                       1000 non-null   int64
9   insured_zip                           1000 non-null   int64
10  insured_sex                           1000 non-null   object
11  insured_education_level               1000 non-null   object
12  insured_occupation                    1000 non-null   object
13  insured_hobbies                       1000 non-null   object
14  insured_relationship                  1000 non-null   object
15  capital-gains                         1000 non-null   int64
16  capital-loss                          1000 non-null   int64
17  incident_date                         1000 non-null   datetime64[ns]
18  incident_type                         1000 non-null   object
19  collision_type                        1000 non-null   object
20  incident_severity                     1000 non-null   object
21  authorities_contacted                 1000 non-null   object
22  incident_state                        1000 non-null   object
23  incident_city                         1000 non-null   object
24  incident_location                     1000 non-null   object
25  incident_hour_of_the_day              1000 non-null   int64
26  number_of_vehicles_involved           1000 non-null   int64
27  property_damage                       1000 non-null   object
28  bodily_injuries                       1000 non-null   int64
29  witnesses                             1000 non-null   int64
30  police_report_available               1000 non-null   object
31  total_claim_amount                    1000 non-null   int64
```

```

32 injury_claim          1000 non-null    int64
33 property_claim        1000 non-null    int64
34 vehicle_claim         1000 non-null    int64
35 auto_make             1000 non-null    object
36 auto_model            1000 non-null    object
37 auto_year             1000 non-null    int64
38 fraud_reported        1000 non-null    object
39 _c39                  0 non-null      float64
dtypes: datetime64[ns](2), float64(2), int64(17), object(19)
memory usage: 312.6+ KB

```

```
[ ]: # Filter the columns with null values
df.isnull().sum()[df.isnull().sum()>0]
```

```
[ ]: _c39      1000
dtype: int64
```

```
[ ]: # Drop _c39 column
df.drop('_c39',axis=1,inplace=True)
```

```
[ ]: # Replace the '?' with nan
df.replace('?',np.nan,inplace=True)
```

```
[ ]: # Total amount of rows
df.shape
```

```
[ ]: (1000, 39)
```

```
[ ]: # Create a list of column names with data type non-numerical
cat_cols=df.select_dtypes(exclude=np.number).columns.tolist()
print(cat_cols)
```

```

['policy_bind_date', 'policy_state', 'policy_csl', 'insured_sex',
'insured_education_level', 'insured_occupation', 'insured_hobbies',
'insured_relationship', 'incident_date', 'incident_type', 'collision_type',
'incident_severity', 'authorities_contacted', 'incident_state', 'incident_city',
'incident_location', 'property_damage', 'police_report_available', 'auto_make',
'auto_model', 'fraud_reported']

```

```
[ ]: # Add the policy number, insurance zip, insurance location to the cat_cols list
cat_cols.extend(['policy_number','insured_zip'])
```

```
[ ]: # Create a function to count the unique values in each column
def unique_values(df):
    for i in df.columns:
        print(i,df[i].nunique())
```

```
[ ]: # Apply the function to the categorical columns  
unique_values(df[cat_cols])
```


```
policy_bind_date 951  
policy_state 3  
policy_csl 3  
insured_sex 2  
insured_education_level 7  
insured_occupation 14  
insured_hobbies 20  
insured_relationship 6  
incident_date 60  
incident_type 4  
collision_type 3  
incident_severity 4  
authorities_contacted 5  
incident_state 7  
incident_city 7  
incident_location 1000  
property_damage 2  
police_report_available 2  
auto_make 14  
auto_model 39  
fraud_reported 2  
policy_number 1000  
insured_zip 995
```

From the set of categorical variables, we can see that some of them have over 900 different values, which is something to keep in mind as we go further into the prediction models.

The variables on this condition are:

- policy\_bind\_date
- incident\_location
- policy\_number
- insured\_zip

```
[ ]: # Let's check the target variable by plotting fraud_reported  
sns.countplot(df['fraud_reported']);
```




```
[ ]: # Count the fraud_reported
df['fraud_reported'].value_counts()
```

```
[ ]: N    753
     Y    247
     Name: fraud_reported, dtype: int64
```

There is a considerable difference between the fraud and non-fraud claims, but is expected given the nature of these events.

```
[ ]: # Histogram of total claim amount
sns.distplot(df['total_claim_amount']);
```



```
[ ]: # Count fraud by Incident State, for Y fraud_reported
df[df['fraud_reported']=='Y']['incident_state'].value_counts()
```

```
[ ]: SC      73
      NY      58
      WV      39
      NC      34
      VA      25
      OH      10
      PA       8
      Name: incident_state, dtype: int64
```

```
[ ]: # Using a USA map to plot the count of fraud_reported = Y by state, color by
      ↪ number of fraud_reported

# Import the required libraries
import plotly.express as px
import plotly.graph_objects as go

# Create a dataframe with the count of fraud_reported = Y by state
df_state=df[df['fraud_reported']=='Y']['incident_state'].value_counts().
      ↪reset_index()
df_state.columns=['state','count']
```



```

# Create a USA map
fig = go.Figure(data=go.Choropleth(
    locations=df_state['state'], # Spatial coordinates
    z = df_state['count'].astype(float), # Data to be color-coded
    locationmode = 'USA-states', # set of locations match entries in `locations`
    colorscale = 'Blues',
    colorbar_title = "Fraud Reported",
))

fig.update_layout(
    title_text = 'Fraud Reported by State',
    geo_scope='usa', # limite map scope to USA
)

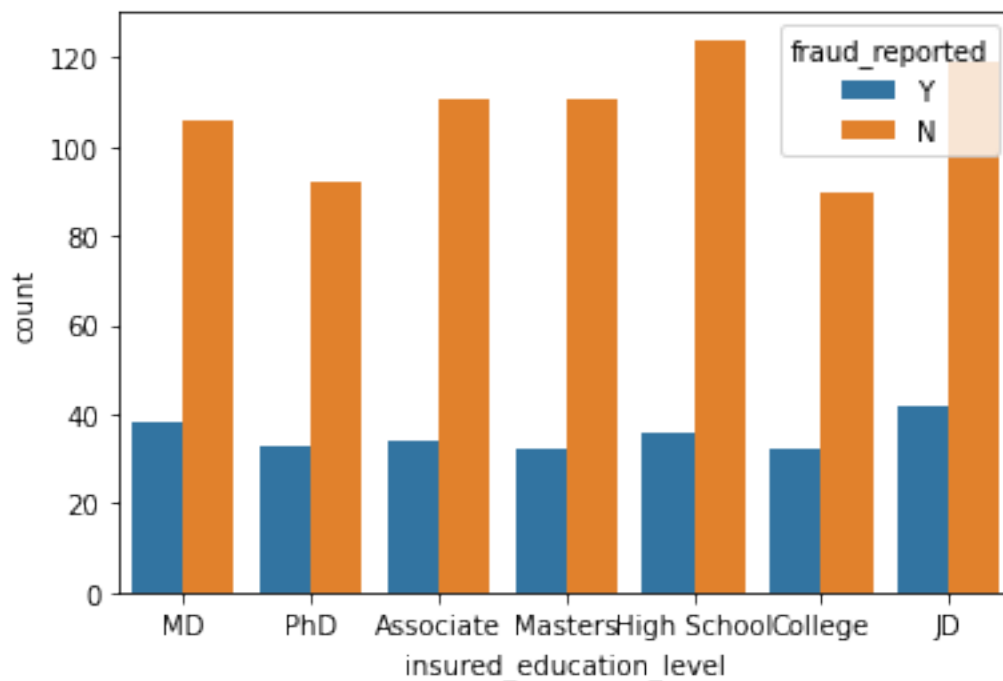
fig.show();

```

```

[ ]: # Plot Breakdown of insured education claim group by fraud_reported
sns.countplot(df['insured_education_level'],hue=df['fraud_reported']);

```

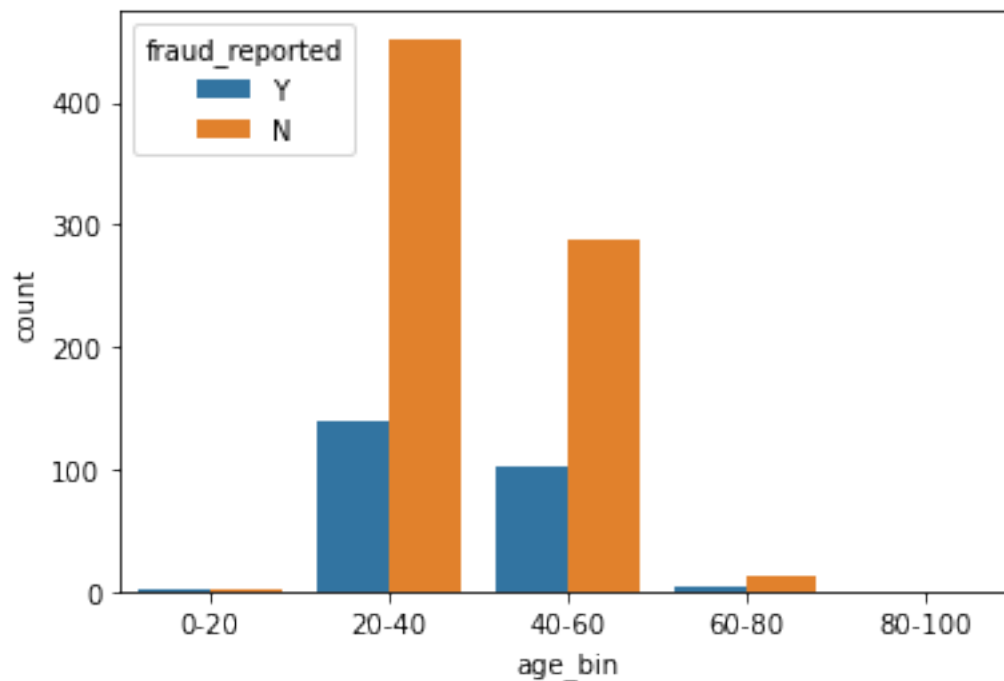


```

[ ]: # Bin the age column group by fraud_reported
df['age_bin']=pd.
    cut(df['age'],bins=[0,20,40,60,80,100],labels=['0-20','20-40','40-60','60-80','80-100'])


```

```
# Plot Breakdown of insured age claim group by fraud_reported
sns.countplot(df['age_bin'],hue=df['fraud_reported']);
```



```
[ ]: # Bin the months_as_customer column group by fraud_reported
df['months_as_customer_bin']=pd.
    ↳cut(df['months_as_customer'],bins=[0,100,200,300,400,500],labels=['0-100','100-200','200-300'])

# Plot Breakdown of insured months_as_customer claim group by fraud_reported
sns.countplot(df['months_as_customer_bin'],hue=df['fraud_reported']);
```



## 4 Data Preprocessing

Considering the for non-numerical variables we will have to generate dummy variables, let's drop the high volume distinct values variables. Let's also add incident date given the is a time series variable.

```
[ ]: cols_to_drop=['policy_number','incident_location','policy_bind_date',
↳ 'insured_zip', 'incident_date', 'age_bin', 'months_as_customer_bin']
```

```
[ ]: # Exclude the cols_to_drop from the cat_cols list
cat_cols=[i for i in cat_cols if i not in cols_to_drop]
```

```
[ ]: # Drop the columns
df.drop(cols_to_drop,axis=1,inplace=True)
```

```
[ ]: # Print the types of the cat_cols
df[cat_cols].dtypes
```

```
[ ]: policy_state      object
policy_csl           object
insured_sex          object
insured_education_level  object
insured_occupation   object
insured_hobbies       object
```

```

insured_relationship    object
incident_type           object
collision_type          object
incident_severity       object
authorities_contacted   object
incident_state          object
incident_city           object
property_damage         object
police_report_available object
auto_make               object
auto_model              object
fraud_reported          object
dtype: object

```

```
[ ]: # Conver the cat_cols to object type
df[cat_cols]=df[cat_cols].astype('string')
```

```
[ ]: # Set the column of cat_cols as factors
for i in cat_cols:
    df[i]=df[i].astype('category')
```

```
[ ]: # Replace the nan values with mode
for i in cat_cols:
    df[i].fillna(df[i].mode()[0],inplace=True)
```

```
[ ]: # Convert the categorical columns to numeric to pass them using string indexer
cat_cols=df.select_dtypes(exclude=np.number).columns.tolist()
cat_cols

# Create a function to convert the categorical columns to numeric
def string_indexer(df,cols):
    for col in cols:
        le=LabelEncoder()
        df[col]=le.fit_transform(df[col])
    return df

# Apply the function to the categorical columns
df_indexed=string_indexer(df,cat_cols)

# Preview the data
df_indexed.head()
```

```
[ ]:   months_as_customer  age  policy_state  policy_csl  policy_deductable  \
0                328    48                2            1            1000
1                228    42                1            1            2000
2                134    29                2            0            2000
3                256    41                0            1            2000
```

4	228	44	0	2	1000
---	-----	----	---	---	------

	policy_annual_premium	umbrella_limit	insured_sex	\
0	1406.91	0	1	
1	1197.22	5000000	1	
2	1413.14	5000000	0	
3	1415.74	6000000	0	
4	1583.91	6000000	1	

	insured_education_level	insured_occupation	...	witnesses	\
0	4	2	...	2	
1	4	6	...	0	
2	6	11	...	3	
3	6	1	...	2	
4	0	11	...	1	

	police_report_available	total_claim_amount	injury_claim	property_claim	\
0	1	71610	6510	13020	
1	0	5070	780	780	
2	0	34650	7700	3850	
3	0	63400	6340	6340	
4	0	6500	1300	650	

	vehicle_claim	auto_make	auto_model	auto_year	fraud_reported
0	52080	10	1	2004	1
1	3510	8	12	2007	1
2	23100	4	30	2007	0
3	50720	3	34	2014	1
4	4550	0	31	2009	0

[5 rows x 34 columns]

## 5 GLM - Gamma Regression for Claim Amount Prediction

The purpose of the following section is to generate a prediction model for the claim amount. The model will be based on a Gamma Regression, which is a generalized linear model (GLM) for predicting continuous positive variables.

As a business problem, insurance companies need to be able to predict the claim amount in order to set the premium for the policy. The premium is the amount of money that the policy holder pays to the insurance company in order to be covered. As other option, the interest to predict claim amount might be rooted on the need to predict the amount of money that the insurance company will have to pay to the policy holder, in a case of a claim, so the company can set aside the resources to react to the claim.

```
[ ]: # Import the sklearn required libraries
from sklearn.linear_model import GammaRegressor
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

```

```

[ ]: # Load the data and split into train and test sets
X = df_indexed.drop('total_claim_amount', axis=1)
y = df_indexed['total_claim_amount']

# Scale the predictor variables
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.4,
    ↪random_state=42)

# Create a gamma regression model
gamma_model = GammaRegressor()

# Fit the model on the training data
gamma_model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = gamma_model.predict(X_test)

```

Mean Squared Error: 116374358.92575014

```

[ ]: # Import the performance metrics library
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

```

```

[ ]: # Calculate the Model Performance Metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
print('MSE:', mean_squared_error(y_test, y_pred))

# Create a function to calculate the adjusted R2
def adj_r2(X,y):
    r2 = gamma_model.score(X,y)
    n = X.shape[0]
    p = X.shape[1]
    adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
    return adjusted_r2

# Calculate the adjusted R2
print('Adjusted R2:', adj_r2(X_test, y_test))

```

R2 Score: 0.8286642792363856

MAE: 9110.823275061093

MSE: 116374358.9257501

Adjusted R2: 0.6368596131291319

The model performance metrics suggest that the gamma regression model is a good fit for the data, with an R2 score of 0.83 indicating that the model explains approximately 83% of the variance in the claim amount.

However, the model's predictions are off by an average of \$9,110.82, as indicated by the mean absolute error (MAE) of 9110.82.

Additionally, the mean squared error (MSE) of 116,374,358.93 suggests that the model's predictions have a larger spread of errors compared to the MAE.

Finally, the adjusted R2 of 0.637 indicates that the model's performance may be slightly impacted by the number of predictor variables used. Overall, the results suggest that the gamma regression model is a good starting point for predicting claim amount, but there may be room for improvement with further model refinement.

## 5.1 Feature Selection

The Gamma model had a considerable amount of variables, for efficiency purposes, and better understanding of the most important variables, we will proceed to generate feature selection through backward elimination.

```
[ ]: # Count the variables in the model
print('Number of variables in the model:', len(gamma_model.coef_))
```

Number of variables in the model: 33

```
[ ]: # Perform a backward feature selection through recursive feature elimination
from sklearn.feature_selection import RFE

# Create the RFE with a gamma regression estimator and 10 features to select
rfe = RFE(estimator=GammaRegressor(), n_features_to_select=10, verbose=1)

# Fit the eliminator to the data
rfe.fit(X_scaled, y)
```

Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.

```

Fitting estimator with 19 features.
Fitting estimator with 18 features.
Fitting estimator with 17 features.
Fitting estimator with 16 features.
Fitting estimator with 15 features.
Fitting estimator with 14 features.
Fitting estimator with 13 features.
Fitting estimator with 12 features.
Fitting estimator with 11 features.

```

```
[ ]: RFE(estimator=GammaRegressor(), n_features_to_select=10, verbose=1)
```

```
[ ]: # Print the features and their ranking (high = dropped early on)
print(dict(zip(X.columns, rfe.ranking_)))
```

```

{'months_as_customer': 16, 'age': 22, 'policy_state': 20, 'policy_csl': 6,
'policy_deductable': 8, 'policy_annual_premium': 7, 'umbrella_limit': 10,
'insured_sex': 14, 'insured_education_level': 2, 'insured_occupation': 24,
'insured_hobbies': 1, 'insured_relationship': 13, 'capital-gains': 12, 'capital-
loss': 17, 'incident_type': 1, 'collision_type': 23, 'incident_severity': 1,
'authorities_contacted': 1, 'incident_state': 3, 'incident_city': 19,
'incident_hour_of_the_day': 1, 'number_of_vehicles_involved': 1,
'property_damage': 21, 'bodily_injuries': 18, 'witnesses': 11,
'police_report_available': 5, 'injury_claim': 1, 'property_claim': 1,
'vehicle_claim': 1, 'auto_make': 9, 'auto_model': 15, 'auto_year': 4,
'fraud_reported': 1}

```

```
[ ]: # Print the features that are not eliminated
print(X.columns[rfe.support_])
```

```

Index(['insured_hobbies', 'incident_type', 'incident_severity',
      'authorities_contacted', 'incident_hour_of_the_day',
      'number_of_vehicles_involved', 'injury_claim', 'property_claim',
      'vehicle_claim', 'fraud_reported'],
      dtype='object')

```

```
[ ]: # Create a dataframe with the features
df_features = pd.DataFrame({'feature':X.columns, 'rank':rfe.ranking_})

# Print the top 10 features
print(df_features.sort_values('rank').head(10))
```

	feature	rank
16	incident_severity	1
28	vehicle_claim	1
27	property_claim	1
26	injury_claim	1
21	number_of_vehicles_involved	1
20	incident_hour_of_the_day	1



17	authorities_contacted	1
14	incident_type	1
10	insured_hobbies	1
32	fraud_reported	1

```
[ ]: # Run the model with the selected features

# Load the data and split into train and test sets

# Select the features from the RFE
X = df_indexed[X.columns[rfe.support_]]
y = df_indexed['total_claim_amount']

# Scale the predictor variables
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.4,
    random_state=42)

# Create a gamma regression model
gamma_model = GammaRegressor()

# Fit the model on the training data
gamma_model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = gamma_model.predict(X_test)
```

```
[ ]: # Calculate the Model Performance Metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
print('MSE:', mean_squared_error(y_test, y_pred))

# Create a function to calculate the adjusted R2
def adj_r2(X,y):
    r2 = gamma_model.score(X,y)
    n = X.shape[0]
    p = X.shape[1]
    adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
    return adjusted_r2

# Calculate the adjusted R2
print('Adjusted R2:', adj_r2(X_test, y_test))
```

R2 Score: 0.8375504064524888

MAE: 8829.271884314403

MSE: 110338738.60385936

Adjusted R2: 0.6640010561580645

The RFE model shows a slight improvement compared to the original model, with a higher R2 score, lower MAE, and lower MSE. Additionally, the adjusted R2 score increased from 0.64 to 0.66, which suggests that the selected features better explain the variation in the response variable.

Overall, the RFE feature selection method was able to identify a subset of features that are more relevant for predicting the total claim amount, resulting in a slightly more accurate model.

## 6 Count Data - Poisson Regression for Amount of Vehicles Involved in the Claim

The purpose of the following section is to generate a prediction model for the amount of vehicles involved in the claims. The model will be based on a Poisson Regression, which is a generalized linear model (GLM) for predicting count data.

By using a Poisson model to predict the number of vehicles involved in an accident claim, the insurance company can estimate the expected number of vehicles involved in a claim and assess the associated risk and potential financial impact. This information can be used to adjust premiums, determine appropriate reserves for future claims, and identify areas for risk mitigation.

Additionally, the insurance company can use the Poisson model to analyze the impact of different variables on the number of vehicles involved in a claim, such as the policyholder's age, sex, education level, occupation, or location. This information can be used to identify high-risk policyholders or regions and develop targeted risk management strategies.

```
[ ]: # Describe the number of vehicle involded  
df['number_of_vehicles_involved'].describe()
```

```
[ ]: count      1000.00000  
      mean       1.83900  
      std       1.01888  
      min       1.00000  
      25%       1.00000  
      50%       1.00000  
      75%       3.00000  
      max       4.00000  
      Name: number_of_vehicles_involved, dtype: float64
```

```
[ ]: from sklearn.linear_model import PoissonRegressor
```

```
[ ]: # Fit a Poisson Model to Predict the Amount of vehicules involved  
  
# Load the data and split into train and test sets  
X = df_indexed.drop('number_of_vehicles_involved', axis=1)  
y = df_indexed['number_of_vehicles_involved']  
  
# Split the data into train and test sets
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.40,
↳random_state=42)

# Scale the predictor variables
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Fit a Poisson regression model on the training data
poisson_model = PoissonRegressor()
poisson_model.fit(X_train_scaled, y_train)

# Make predictions on the test data
y_pred = poisson_model.predict(X_test_scaled)

```

```

[ ]: # Evaluate the model performance
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
print('MSE:', mean_squared_error(y_test, y_pred))

# Create a function to calculate the adjusted R2
def adj_r2(X,y):
    r2 = poisson_model.score(X,y)
    n = X.shape[0]
    p = X.shape[1]
    adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
    return adjusted_r2

# Calculate the adjusted R2
print('Adjusted R2:', adj_r2(X_test_scaled, y_test))

```

R2 Score: 0.7227489970960406

MAE: 0.47486946204652375

MSE: 0.2938427426089932

Adjusted R2: 0.7168411325544817

The R2 score of 0.72 indicates that the Poisson model has a moderate level of fit and is able to explain around 72% of the variance in the number of vehicles involved in a claim.

The MAE value of 0.47 suggests that the model is able to predict the number of vehicles involved with an average error of 0.47, which may or may not be acceptable depending on the context of the problem.

The MSE value of 0.29 suggests that the model's predictions have a moderate level of variance, which means that there is some amount of variability in the predictions around the true values.

The adjusted R2 value of 0.71 indicates that the model has a good level of fit, taking into account the number of predictors used in the model. Overall, the model seems to have a moderate level of fit, but it may benefit from further refinement or the inclusion of additional predictors.

```
[ ]: # Predict the number of vehicles involved in an accident for a random
      ↪ observation
print('Predicted number of vehicles involved:', poisson_model.
      ↪ predict(X_test_scaled[0].reshape(1,-1))[0])
print('Actual number of vehicles involved:', y_test.iloc[0])
```

Predicted number of vehicles involved: 1.4452792352608677  
 Actual number of vehicles involved: 1

```
[ ]: # Calculate the probability of accidents involving 1, 2, 3, and 4 vehicles

from scipy.stats import poisson

# Create a function to calculate the probability of accidents involving 1, 2,
↪ 3, and 4 vehicles given the model, and returning a single percentage per
↪ number of vehicles
def prob_vehicles(model, X, num_vehicles):
    # Convert X to a DataFrame
    X_df = pd.DataFrame(X, columns=X_test.columns)
    # Create a dataframe with the features
    df_features = pd.DataFrame({'feature':X_df.columns, 'coef':model.coef_})
    # Create a dataframe with the features and their exponential
    df_features['exp_coef'] = df_features['coef'].apply(lambda x: np.exp(x))
    # Create a dataframe with the features and their exponential multiplied by
    ↪ the number of vehicles
    df_features['exp_coef_num_vehicles'] = df_features['exp_coef'].apply(lambda
    ↪ x: x**num_vehicles)
    # Calculate the probability of accidents involving 1, 2, 3, and 4 vehicles
    prob = df_features['exp_coef_num_vehicles'].prod()
    return prob

# Calculate the probability of accidents involving 1, 2, 3, and 4 vehicles
print('Probability of accidents involving 1 vehicle:',
      ↪ prob_vehicles(poisson_model, X_test_scaled[0].reshape(1,-1), 1))
print('Probability of accidents involving 2 vehicles:',
      ↪ prob_vehicles(poisson_model, X_test_scaled[0].reshape(1,-1), 2))
print('Probability of accidents involving 3 vehicles:',
      ↪ prob_vehicles(poisson_model, X_test_scaled[0].reshape(1,-1), 3))
print('Probability of accidents involving 4 vehicles:',
      ↪ prob_vehicles(poisson_model, X_test_scaled[0].reshape(1,-1), 4))
```

Probability of accidents involving 1 vehicle: 0.7766158989555773  
 Probability of accidents involving 2 vehicles: 0.6031322545105795  
 Probability of accidents involving 3 vehicles: 0.46840209802583743  
 Probability of accidents involving 4 vehicles: 0.36376851643101427

The results suggest that the probability of an accident involving 1 vehicle is the highest among the four options, with a probability of approximately 0.78. The probability decreases as the number of

vehicles involved in the accident increases, with a probability of approximately 0.60 for accidents involving 2 vehicles, approximately 0.47 for accidents involving 3 vehicles, and approximately 0.36 for accidents involving 4 vehicles.

These probabilities may provide insights for insurance companies and policy makers when assessing risk and designing policies related to auto accidents. It may be helpful to investigate the causes of accidents involving multiple vehicles and to determine ways to mitigate these risks, such as improving road infrastructure and safety measures.

## 7 Neural Network for Fraud Prediction

The following section is focused on generating a prediction model for the fraud probability. The model will be based on a Neural Network Model, which is a machine learning model that can be used to predict the probability of an event occurring.

As a business problem, a machine learning model can help insurance companies to detect and prevent fraud more accurately and efficiently, reducing financial losses and maintaining lower premiums for policyholders. By analyzing historical data and identifying patterns indicative of fraud, the model can prioritize investigations and allocate resources more effectively, improving overall risk management for the company.

```
[ ]: # Define the dependent and independent variables
X = df.drop('fraud_reported', axis=1)
y = df['fraud_reported']

[ ]: # Drop fraud_reported from the cat_cols list
cat_cols=[i for i in cat_cols if i not in ['fraud_reported']]

[ ]: # Get the dummy variables for the categorical columns
X=pd.get_dummies(X,columns=cat_cols,drop_first=True)

[ ]: # Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5,
↳random_state=52)

[ ]: # Fit a neural network model to predict fraud_reported using Keras

# Import the required libraries
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping

# Create a function to create a neural network model
def create_model():
    # create model
    model = Sequential()
    model.add(Dense(16, input_dim=X_train.shape[1], activation='relu'))
```

```

model.add(Dense(8, activation='relu'))
model.add(Dense(4, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

# Compile model
adam=Adam(lr=0.0001)
model.compile(loss='binary_crossentropy', optimizer=adam,
metrics=['accuracy'])
return model

# Create a model
model = create_model()

# Print the model summary
print(model.summary())

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	1584
dense_1 (Dense)	(None, 8)	136
dense_2 (Dense)	(None, 4)	36
dense_3 (Dense)	(None, 1)	5

Total params: 1,761  
 Trainable params: 1,761  
 Non-trainable params: 0

None

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	1584
dense_1 (Dense)	(None, 8)	136
dense_2 (Dense)	(None, 4)	36
dense_3 (Dense)	(None, 1)	5

Total params: 1,761  
 Trainable params: 1,761

Non-trainable params: 0

-----  
None

```
[ ]: # Fit the model  
model.fit(X_train, y_train, epochs=200, batch_size=20, verbose=1)
```

```
Epoch 1/200  
25/25 [=====] - 0s 899us/step - loss: 6947.1743 -  
accuracy: 0.6580  
Epoch 2/200  
25/25 [=====] - 0s 898us/step - loss: 2474.5608 -  
accuracy: 0.6360  
Epoch 3/200  
25/25 [=====] - 0s 965us/step - loss: 908.8033 -  
accuracy: 0.5800  
Epoch 4/200  
25/25 [=====] - 0s 792us/step - loss: 745.7484 -  
accuracy: 0.5840  
Epoch 5/200  
25/25 [=====] - 0s 815us/step - loss: 552.3362 -  
accuracy: 0.5800  
Epoch 6/200  
25/25 [=====] - 0s 836us/step - loss: 541.5812 -  
accuracy: 0.5740  
Epoch 7/200  
25/25 [=====] - 0s 878us/step - loss: 586.6372 -  
accuracy: 0.5820  
Epoch 8/200  
25/25 [=====] - 0s 881us/step - loss: 542.7558 -  
accuracy: 0.5820  
Epoch 9/200  
25/25 [=====] - 0s 819us/step - loss: 456.3873 -  
accuracy: 0.6080  
Epoch 10/200  
25/25 [=====] - 0s 837us/step - loss: 492.1421 -  
accuracy: 0.6000  
Epoch 11/200  
25/25 [=====] - 0s 836us/step - loss: 519.8999 -  
accuracy: 0.6120  
Epoch 12/200  
25/25 [=====] - 0s 794us/step - loss: 472.0103 -  
accuracy: 0.6100  
Epoch 13/200  
25/25 [=====] - 0s 877us/step - loss: 398.6445 -  
accuracy: 0.6040  
Epoch 14/200  
25/25 [=====] - 0s 836us/step - loss: 380.9908 -
```

```

accuracy: 0.6020
Epoch 15/200
25/25 [=====] - 0s 839us/step - loss: 392.1096 -
accuracy: 0.6140
Epoch 16/200
25/25 [=====] - 0s 814us/step - loss: 398.7802 -
accuracy: 0.6080
Epoch 17/200
25/25 [=====] - 0s 1ms/step - loss: 381.8946 -
accuracy: 0.6140
Epoch 18/200
25/25 [=====] - 0s 1ms/step - loss: 367.1260 -
accuracy: 0.5940
Epoch 19/200
25/25 [=====] - 0s 1ms/step - loss: 336.5479 -
accuracy: 0.6120
Epoch 20/200
25/25 [=====] - 0s 943us/step - loss: 402.0947 -
accuracy: 0.6280
Epoch 21/200
25/25 [=====] - 0s 877us/step - loss: 322.5425 -
accuracy: 0.6160
Epoch 22/200
25/25 [=====] - 0s 814us/step - loss: 343.7301 -
accuracy: 0.6160
Epoch 23/200
25/25 [=====] - 0s 879us/step - loss: 316.4546 -
accuracy: 0.6160
Epoch 24/200
25/25 [=====] - 0s 835us/step - loss: 318.9844 -
accuracy: 0.6160
Epoch 25/200
25/25 [=====] - 0s 826us/step - loss: 250.8297 -
accuracy: 0.6200
Epoch 26/200
25/25 [=====] - 0s 818us/step - loss: 267.4481 -
accuracy: 0.6020
Epoch 27/200
25/25 [=====] - 0s 814us/step - loss: 336.0552 -
accuracy: 0.6280
Epoch 28/200
25/25 [=====] - 0s 857us/step - loss: 250.4970 -
accuracy: 0.6120
Epoch 29/200
25/25 [=====] - 0s 835us/step - loss: 280.1104 -
accuracy: 0.6100
Epoch 30/200
25/25 [=====] - 0s 836us/step - loss: 237.4940 -

```



```

accuracy: 0.6260
Epoch 31/200
25/25 [=====] - 0s 815us/step - loss: 284.4987 -
accuracy: 0.5940
Epoch 32/200
25/25 [=====] - 0s 839us/step - loss: 399.9630 -
accuracy: 0.6040
Epoch 33/200
25/25 [=====] - 0s 846us/step - loss: 280.7903 -
accuracy: 0.6120
Epoch 34/200
25/25 [=====] - 0s 1ms/step - loss: 200.7180 -
accuracy: 0.5940
Epoch 35/200
25/25 [=====] - 0s 981us/step - loss: 267.8877 -
accuracy: 0.6000
Epoch 36/200
25/25 [=====] - 0s 1ms/step - loss: 236.3749 -
accuracy: 0.5960
Epoch 37/200
25/25 [=====] - 0s 865us/step - loss: 168.1265 -
accuracy: 0.5800
Epoch 38/200
25/25 [=====] - 0s 841us/step - loss: 202.3585 -
accuracy: 0.5640
Epoch 39/200
25/25 [=====] - 0s 876us/step - loss: 141.2833 -
accuracy: 0.5520
Epoch 40/200
25/25 [=====] - 0s 835us/step - loss: 143.3056 -
accuracy: 0.5480
Epoch 41/200
25/25 [=====] - 0s 898us/step - loss: 146.7035 -
accuracy: 0.5360
Epoch 42/200
25/25 [=====] - 0s 815us/step - loss: 163.3881 -
accuracy: 0.5700
Epoch 43/200
25/25 [=====] - 0s 796us/step - loss: 143.7869 -
accuracy: 0.5840
Epoch 44/200
25/25 [=====] - 0s 839us/step - loss: 247.4911 -
accuracy: 0.5560
Epoch 45/200
25/25 [=====] - 0s 877us/step - loss: 208.7947 -
accuracy: 0.5760
Epoch 46/200
25/25 [=====] - 0s 856us/step - loss: 126.1005 -

```

```

accuracy: 0.6100
Epoch 47/200
25/25 [=====] - 0s 793us/step - loss: 141.9734 -
accuracy: 0.6080
Epoch 48/200
25/25 [=====] - 0s 983us/step - loss: 144.8161 -
accuracy: 0.6200
Epoch 49/200
25/25 [=====] - 0s 1ms/step - loss: 125.4669 -
accuracy: 0.6080
Epoch 50/200
25/25 [=====] - 0s 945us/step - loss: 152.8707 -
accuracy: 0.6160
Epoch 51/200
25/25 [=====] - 0s 814us/step - loss: 129.0500 -
accuracy: 0.6420
Epoch 52/200
25/25 [=====] - 0s 814us/step - loss: 72.6257 -
accuracy: 0.6340
Epoch 53/200
25/25 [=====] - 0s 835us/step - loss: 98.1095 -
accuracy: 0.6400
Epoch 54/200
25/25 [=====] - 0s 796us/step - loss: 84.0623 -
accuracy: 0.6340
Epoch 55/200
25/25 [=====] - 0s 815us/step - loss: 56.1537 -
accuracy: 0.6460
Epoch 56/200
25/25 [=====] - 0s 796us/step - loss: 91.9431 -
accuracy: 0.6440
Epoch 57/200
25/25 [=====] - 0s 814us/step - loss: 81.1964 -
accuracy: 0.6440
Epoch 58/200
25/25 [=====] - 0s 814us/step - loss: 101.5889 -
accuracy: 0.6460
Epoch 59/200
25/25 [=====] - 0s 814us/step - loss: 91.1583 -
accuracy: 0.6640
Epoch 60/200
25/25 [=====] - 0s 794us/step - loss: 55.9413 -
accuracy: 0.6360
Epoch 61/200
25/25 [=====] - 0s 797us/step - loss: 139.1858 -
accuracy: 0.6600
Epoch 62/200
25/25 [=====] - 0s 796us/step - loss: 82.1207 -

```

```

accuracy: 0.6520
Epoch 63/200
25/25 [=====] - 0s 982us/step - loss: 103.8644 -
accuracy: 0.6680
Epoch 64/200
25/25 [=====] - 0s 1ms/step - loss: 134.1967 -
accuracy: 0.6680
Epoch 65/200
25/25 [=====] - 0s 981us/step - loss: 53.8054 -
accuracy: 0.6740
Epoch 66/200
25/25 [=====] - 0s 794us/step - loss: 54.0155 -
accuracy: 0.6880
Epoch 67/200
25/25 [=====] - 0s 798us/step - loss: 193.8517 -
accuracy: 0.6700
Epoch 68/200
25/25 [=====] - 0s 794us/step - loss: 60.0211 -
accuracy: 0.6900
Epoch 69/200
25/25 [=====] - 0s 814us/step - loss: 77.5990 -
accuracy: 0.6780
Epoch 70/200
25/25 [=====] - 0s 835us/step - loss: 55.0239 -
accuracy: 0.6660
Epoch 71/200
25/25 [=====] - 0s 791us/step - loss: 62.3729 -
accuracy: 0.6720
Epoch 72/200
25/25 [=====] - 0s 797us/step - loss: 104.9686 -
accuracy: 0.6620
Epoch 73/200
25/25 [=====] - 0s 796us/step - loss: 68.0541 -
accuracy: 0.6760
Epoch 74/200
25/25 [=====] - 0s 775us/step - loss: 165.1224 -
accuracy: 0.6880
Epoch 75/200
25/25 [=====] - 0s 796us/step - loss: 86.6320 -
accuracy: 0.6660
Epoch 76/200
25/25 [=====] - 0s 773us/step - loss: 101.0446 -
accuracy: 0.6580
Epoch 77/200
25/25 [=====] - 0s 817us/step - loss: 55.4335 -
accuracy: 0.6880
Epoch 78/200
25/25 [=====] - 0s 795us/step - loss: 56.9533 -

```

```

accuracy: 0.6880
Epoch 79/200
25/25 [=====] - 0s 1ms/step - loss: 101.0332 -
accuracy: 0.6740
Epoch 80/200
25/25 [=====] - 0s 1ms/step - loss: 100.4065 -
accuracy: 0.6820
Epoch 81/200
25/25 [=====] - 0s 1ms/step - loss: 102.4389 -
accuracy: 0.6680
Epoch 82/200
25/25 [=====] - 0s 1ms/step - loss: 82.6608 - accuracy:
0.6920
Epoch 83/200
25/25 [=====] - 0s 836us/step - loss: 84.9176 -
accuracy: 0.6760
Epoch 84/200
25/25 [=====] - 0s 836us/step - loss: 84.8463 -
accuracy: 0.6840
Epoch 85/200
25/25 [=====] - 0s 800us/step - loss: 65.7322 -
accuracy: 0.6760
Epoch 86/200
25/25 [=====] - 0s 794us/step - loss: 58.2357 -
accuracy: 0.6760
Epoch 87/200
25/25 [=====] - 0s 793us/step - loss: 96.6498 -
accuracy: 0.6780
Epoch 88/200
25/25 [=====] - 0s 856us/step - loss: 121.9459 -
accuracy: 0.6800
Epoch 89/200
25/25 [=====] - 0s 841us/step - loss: 55.6262 -
accuracy: 0.6780
Epoch 90/200
25/25 [=====] - 0s 793us/step - loss: 68.2254 -
accuracy: 0.6600
Epoch 91/200
25/25 [=====] - 0s 835us/step - loss: 53.3755 -
accuracy: 0.6880
Epoch 92/200
25/25 [=====] - 0s 814us/step - loss: 46.2688 -
accuracy: 0.6880
Epoch 93/200
25/25 [=====] - 0s 795us/step - loss: 46.6621 -
accuracy: 0.6900
Epoch 94/200
25/25 [=====] - 0s 962us/step - loss: 57.2078 -

```

```

accuracy: 0.6860
Epoch 95/200
25/25 [=====] - 0s 1ms/step - loss: 49.6593 - accuracy:
0.6960
Epoch 96/200
25/25 [=====] - 0s 905us/step - loss: 80.0890 -
accuracy: 0.6880
Epoch 97/200
25/25 [=====] - 0s 794us/step - loss: 70.6270 -
accuracy: 0.6700
Epoch 98/200
25/25 [=====] - 0s 835us/step - loss: 69.6082 -
accuracy: 0.6840
Epoch 99/200
25/25 [=====] - 0s 794us/step - loss: 67.1859 -
accuracy: 0.6840
Epoch 100/200
25/25 [=====] - 0s 773us/step - loss: 64.1913 -
accuracy: 0.6800
Epoch 101/200
25/25 [=====] - 0s 793us/step - loss: 71.5268 -
accuracy: 0.6800
Epoch 102/200
25/25 [=====] - 0s 772us/step - loss: 71.3787 -
accuracy: 0.6920
Epoch 103/200
25/25 [=====] - 0s 815us/step - loss: 50.0667 -
accuracy: 0.6840
Epoch 104/200
25/25 [=====] - 0s 815us/step - loss: 58.8049 -
accuracy: 0.6880
Epoch 105/200
25/25 [=====] - 0s 795us/step - loss: 64.7363 -
accuracy: 0.6640
Epoch 106/200
25/25 [=====] - 0s 794us/step - loss: 88.3265 -
accuracy: 0.7020
Epoch 107/200
25/25 [=====] - 0s 835us/step - loss: 97.6908 -
accuracy: 0.7000
Epoch 108/200
25/25 [=====] - 0s 772us/step - loss: 54.4678 -
accuracy: 0.6980
Epoch 109/200
25/25 [=====] - 0s 793us/step - loss: 58.5719 -
accuracy: 0.6960
Epoch 110/200
25/25 [=====] - 0s 1ms/step - loss: 69.8249 - accuracy:

```

```

0.7120
Epoch 111/200
25/25 [=====] - 0s 1ms/step - loss: 86.3784 - accuracy:
0.7080
Epoch 112/200
25/25 [=====] - 0s 1ms/step - loss: 42.0706 - accuracy:
0.7000
Epoch 113/200
25/25 [=====] - 0s 794us/step - loss: 35.6461 -
accuracy: 0.7080
Epoch 114/200
25/25 [=====] - 0s 793us/step - loss: 44.5446 -
accuracy: 0.6740
Epoch 115/200
25/25 [=====] - 0s 793us/step - loss: 53.1757 -
accuracy: 0.6720
Epoch 116/200
25/25 [=====] - 0s 793us/step - loss: 31.8403 -
accuracy: 0.7100
Epoch 117/200
25/25 [=====] - 0s 794us/step - loss: 63.2532 -
accuracy: 0.6660
Epoch 118/200
25/25 [=====] - 0s 799us/step - loss: 51.0646 -
accuracy: 0.6920
Epoch 119/200
25/25 [=====] - 0s 818us/step - loss: 159.9407 -
accuracy: 0.6780
Epoch 120/200
25/25 [=====] - 0s 792us/step - loss: 102.6931 -
accuracy: 0.6980
Epoch 121/200
25/25 [=====] - 0s 834us/step - loss: 26.7488 -
accuracy: 0.7080
Epoch 122/200
25/25 [=====] - 0s 794us/step - loss: 48.9884 -
accuracy: 0.6980
Epoch 123/200
25/25 [=====] - 0s 794us/step - loss: 69.0593 -
accuracy: 0.6960
Epoch 124/200
25/25 [=====] - 0s 2ms/step - loss: 80.7219 - accuracy:
0.6860
Epoch 125/200
25/25 [=====] - 0s 954us/step - loss: 28.2026 -
accuracy: 0.7060
Epoch 126/200
25/25 [=====] - 0s 882us/step - loss: 34.3212 -

```

```

accuracy: 0.7160
Epoch 127/200
25/25 [=====] - 0s 813us/step - loss: 43.6475 -
accuracy: 0.6760
Epoch 128/200
25/25 [=====] - 0s 814us/step - loss: 86.1366 -
accuracy: 0.6980
Epoch 129/200
25/25 [=====] - 0s 776us/step - loss: 106.6821 -
accuracy: 0.6880
Epoch 130/200
25/25 [=====] - 0s 816us/step - loss: 171.6767 -
accuracy: 0.6700
Epoch 131/200
25/25 [=====] - 0s 798us/step - loss: 114.7948 -
accuracy: 0.7080
Epoch 132/200
25/25 [=====] - 0s 793us/step - loss: 130.5965 -
accuracy: 0.6880
Epoch 133/200
25/25 [=====] - 0s 772us/step - loss: 39.1164 -
accuracy: 0.7100
Epoch 134/200
25/25 [=====] - 0s 792us/step - loss: 57.6070 -
accuracy: 0.7100
Epoch 135/200
25/25 [=====] - 0s 793us/step - loss: 37.7200 -
accuracy: 0.7060
Epoch 136/200
25/25 [=====] - 0s 815us/step - loss: 57.3084 -
accuracy: 0.7060
Epoch 137/200
25/25 [=====] - 0s 1ms/step - loss: 79.1852 - accuracy:
0.7040
Epoch 138/200
25/25 [=====] - 0s 1ms/step - loss: 76.8136 - accuracy:
0.6980
Epoch 139/200
25/25 [=====] - 0s 960us/step - loss: 74.8812 -
accuracy: 0.7060
Epoch 140/200
25/25 [=====] - 0s 814us/step - loss: 51.2769 -
accuracy: 0.7160
Epoch 141/200
25/25 [=====] - 0s 815us/step - loss: 123.5782 -
accuracy: 0.6940
Epoch 142/200
25/25 [=====] - 0s 772us/step - loss: 97.9099 -

```

```

accuracy: 0.6900
Epoch 143/200
25/25 [=====] - 0s 813us/step - loss: 100.2137 -
accuracy: 0.7260
Epoch 144/200
25/25 [=====] - 0s 818us/step - loss: 163.4574 -
accuracy: 0.6940
Epoch 145/200
25/25 [=====] - 0s 836us/step - loss: 262.9234 -
accuracy: 0.6940
Epoch 146/200
25/25 [=====] - 0s 989us/step - loss: 190.6122 -
accuracy: 0.7120
Epoch 147/200
25/25 [=====] - 0s 793us/step - loss: 50.9315 -
accuracy: 0.7060
Epoch 148/200
25/25 [=====] - 0s 792us/step - loss: 93.2805 -
accuracy: 0.6980
Epoch 149/200
25/25 [=====] - 0s 794us/step - loss: 61.3404 -
accuracy: 0.7280
Epoch 150/200
25/25 [=====] - 0s 753us/step - loss: 47.1367 -
accuracy: 0.6960
Epoch 151/200
25/25 [=====] - 0s 838us/step - loss: 128.6801 -
accuracy: 0.7220
Epoch 152/200
25/25 [=====] - 0s 1ms/step - loss: 103.4937 -
accuracy: 0.7220
Epoch 153/200
25/25 [=====] - 0s 1ms/step - loss: 72.6324 - accuracy:
0.7020
Epoch 154/200
25/25 [=====] - 0s 920us/step - loss: 53.5139 -
accuracy: 0.7120
Epoch 155/200
25/25 [=====] - 0s 773us/step - loss: 69.2190 -
accuracy: 0.7080
Epoch 156/200
25/25 [=====] - 0s 772us/step - loss: 79.1296 -
accuracy: 0.7080
Epoch 157/200
25/25 [=====] - 0s 781us/step - loss: 48.9102 -
accuracy: 0.7020
Epoch 158/200
25/25 [=====] - 0s 817us/step - loss: 31.8400 -

```



```

accuracy: 0.7260
Epoch 159/200
25/25 [=====] - 0s 756us/step - loss: 30.7672 -
accuracy: 0.7280
Epoch 160/200
25/25 [=====] - 0s 813us/step - loss: 57.0757 -
accuracy: 0.7260
Epoch 161/200
25/25 [=====] - 0s 836us/step - loss: 86.2038 -
accuracy: 0.7220
Epoch 162/200
25/25 [=====] - 0s 794us/step - loss: 44.8307 -
accuracy: 0.7140
Epoch 163/200
25/25 [=====] - 0s 794us/step - loss: 29.2588 -
accuracy: 0.7300
Epoch 164/200
25/25 [=====] - 0s 773us/step - loss: 54.9562 -
accuracy: 0.7200
Epoch 165/200
25/25 [=====] - 0s 796us/step - loss: 49.0006 -
accuracy: 0.7300
Epoch 166/200
25/25 [=====] - 0s 1ms/step - loss: 64.4708 - accuracy:
0.7020
Epoch 167/200
25/25 [=====] - 0s 983us/step - loss: 23.1144 -
accuracy: 0.7280
Epoch 168/200
25/25 [=====] - 0s 860us/step - loss: 55.4248 -
accuracy: 0.7280
Epoch 169/200
25/25 [=====] - 0s 798us/step - loss: 49.9142 -
accuracy: 0.7220
Epoch 170/200
25/25 [=====] - 0s 793us/step - loss: 83.9214 -
accuracy: 0.7120
Epoch 171/200
25/25 [=====] - 0s 794us/step - loss: 107.3120 -
accuracy: 0.7240
Epoch 172/200
25/25 [=====] - 0s 795us/step - loss: 40.4449 -
accuracy: 0.7240
Epoch 173/200
25/25 [=====] - 0s 795us/step - loss: 64.2503 -
accuracy: 0.7020
Epoch 174/200
25/25 [=====] - 0s 781us/step - loss: 144.2501 -

```

```

accuracy: 0.7080
Epoch 175/200
25/25 [=====] - 0s 751us/step - loss: 46.9898 -
accuracy: 0.7220
Epoch 176/200
25/25 [=====] - 0s 837us/step - loss: 21.0693 -
accuracy: 0.7380
Epoch 177/200
25/25 [=====] - 0s 796us/step - loss: 41.0686 -
accuracy: 0.7260
Epoch 178/200
25/25 [=====] - 0s 752us/step - loss: 69.5288 -
accuracy: 0.7120
Epoch 179/200
25/25 [=====] - 0s 877us/step - loss: 91.4239 -
accuracy: 0.7020
Epoch 180/200
25/25 [=====] - 0s 2ms/step - loss: 177.0196 -
accuracy: 0.7360
Epoch 181/200
25/25 [=====] - 0s 965us/step - loss: 114.6377 -
accuracy: 0.7100
Epoch 182/200
25/25 [=====] - 0s 881us/step - loss: 103.1633 -
accuracy: 0.7140
Epoch 183/200
25/25 [=====] - 0s 796us/step - loss: 37.5359 -
accuracy: 0.7360
Epoch 184/200
25/25 [=====] - 0s 815us/step - loss: 38.0279 -
accuracy: 0.7200
Epoch 185/200
25/25 [=====] - 0s 794us/step - loss: 47.0279 -
accuracy: 0.7180
Epoch 186/200
25/25 [=====] - 0s 836us/step - loss: 52.1423 -
accuracy: 0.7260
Epoch 187/200
25/25 [=====] - 0s 779us/step - loss: 36.5526 -
accuracy: 0.7240
Epoch 188/200
25/25 [=====] - 0s 835us/step - loss: 41.8313 -
accuracy: 0.7340
Epoch 189/200
25/25 [=====] - 0s 818us/step - loss: 47.0684 -
accuracy: 0.7220
Epoch 190/200
25/25 [=====] - 0s 793us/step - loss: 28.6111 -

```

```

accuracy: 0.7300
Epoch 191/200
25/25 [=====] - 0s 800us/step - loss: 59.4901 -
accuracy: 0.7340
Epoch 192/200
25/25 [=====] - 0s 794us/step - loss: 95.7343 -
accuracy: 0.7220
Epoch 193/200
25/25 [=====] - 0s 941us/step - loss: 84.1668 -
accuracy: 0.7180
Epoch 194/200
25/25 [=====] - 0s 1ms/step - loss: 34.4738 - accuracy:
0.7420
Epoch 195/200
25/25 [=====] - 0s 901us/step - loss: 128.3673 -
accuracy: 0.7360
Epoch 196/200
25/25 [=====] - 0s 814us/step - loss: 141.1569 -
accuracy: 0.7020
Epoch 197/200
25/25 [=====] - 0s 773us/step - loss: 87.6433 -
accuracy: 0.7300
Epoch 198/200
25/25 [=====] - 0s 796us/step - loss: 42.1844 -
accuracy: 0.7200
Epoch 199/200
25/25 [=====] - 0s 836us/step - loss: 34.2162 -
accuracy: 0.7240
Epoch 200/200
25/25 [=====] - 0s 798us/step - loss: 58.8181 -
accuracy: 0.7280

```

```
[ ]: <keras.callbacks.History at 0x145d278d550>
```

```
[ ]: # Predict the model
y_pred = model.predict(X_test)
```

```

16/16 [=====] - 0s 669us/step
16/16 [=====] - 0s 669us/step

```

```
[ ]: # Print the accuracy score
print(accuracy_score(y_test,y_pred.round()))
```


```
0.728
```

```
[ ]: # Print the confusion matrix
print(confusion_matrix(y_test,y_pred.round()))
```

```
[[361  20]
```

```
[116 3]]
```

```
[ ]: # Plot the confusion matrix
sns.heatmap(confusion_matrix(y_test,y_pred.round()),annot=True,fmt='d');
```



The model accuracy does not seem that great, with only a score of 73% accuracy. Let's see if we can improve it by hyperparameter tuning.

```
[ ]: # Import required libraries
import numpy as np
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam

# Define a function to create a Keras model
def create_model(learning_rate=0.01, activation='relu'):
    # create model
    model = Sequential()
    model.add(Dense(16, input_dim=X_train.shape[1], activation=activation))
    model.add(Dense(8, activation=activation))
    model.add(Dense(4, activation=activation))
    model.add(Dense(1, activation='sigmoid'))

    # Compile model
```

```

    optimizer = Adam(lr=learning_rate)
    model.compile(loss='binary_crossentropy', optimizer=optimizer,
metrics=['accuracy'])
    return model

# Define hyperparameters
param_grid = {'batch_size': [20, 40, 60, 80, 100],
              'epochs': [100, 200, 300, 400, 500],
              'learning_rate': [0.01, 0.001, 0.0001],
              'activation': ['relu', 'tanh']}

# Create a KerasClassifier object
model = KerasClassifier(build_fn=create_model)

# Create a GridSearchCV object
grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=3)

# Fit the GridSearchCV object with the data
grid_result = grid.fit(X_train, y_train)

# Print the best score and best parameters
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))

```

```

Epoch 1/100
7/7 [=====] - 1s 1ms/step - loss: 38319.8867 -
accuracy: 0.4680
Epoch 2/100
7/7 [=====] - 0s 1ms/step - loss: 25369.8164 -
accuracy: 0.5080
Epoch 3/100
7/7 [=====] - 0s 1ms/step - loss: 11500.5107 -
accuracy: 0.6280
Epoch 4/100
7/7 [=====] - 0s 1ms/step - loss: 1643.5712 - accuracy:
0.7040
Epoch 5/100
7/7 [=====] - 0s 1ms/step - loss: 2228.5325 - accuracy:
0.7300
Epoch 6/100
7/7 [=====] - 0s 1ms/step - loss: 1689.7468 - accuracy:
0.7040
Epoch 7/100
7/7 [=====] - 0s 1ms/step - loss: 500.6108 - accuracy:
0.5720
Epoch 8/100
7/7 [=====] - 0s 1ms/step - loss: 483.5487 - accuracy:
0.6080
Epoch 9/100

```

7/7 [=====] - 0s 1ms/step - loss: 464.6248 - accuracy:  
 0.6800  
 Epoch 10/100  
 7/7 [=====] - 0s 1ms/step - loss: 476.2349 - accuracy:  
 0.6120  
 Epoch 11/100  
 7/7 [=====] - 0s 1ms/step - loss: 337.5928 - accuracy:  
 0.6680  
 Epoch 12/100  
 7/7 [=====] - 0s 1ms/step - loss: 271.8499 - accuracy:  
 0.6240  
 Epoch 13/100  
 7/7 [=====] - 0s 1ms/step - loss: 182.0222 - accuracy:  
 0.6640  
 Epoch 14/100  
 7/7 [=====] - 0s 2ms/step - loss: 173.7761 - accuracy:  
 0.6400  
 Epoch 15/100  
 7/7 [=====] - 0s 2ms/step - loss: 143.0125 - accuracy:  
 0.6160  
 Epoch 16/100  
 7/7 [=====] - 0s 2ms/step - loss: 214.9517 - accuracy:  
 0.6540  
 Epoch 17/100  
 7/7 [=====] - 0s 1ms/step - loss: 162.1337 - accuracy:  
 0.6360  
 Epoch 18/100  
 7/7 [=====] - 0s 1ms/step - loss: 130.4715 - accuracy:  
 0.6560  
 Epoch 19/100  
 7/7 [=====] - 0s 1ms/step - loss: 182.2525 - accuracy:  
 0.6420  
 Epoch 20/100  
 7/7 [=====] - 0s 1ms/step - loss: 189.3147 - accuracy:  
 0.6520  
 Epoch 21/100  
 7/7 [=====] - 0s 1ms/step - loss: 171.3895 - accuracy:  
 0.6760  
 Epoch 22/100  
 7/7 [=====] - 0s 2ms/step - loss: 184.2287 - accuracy:  
 0.6460  
 Epoch 23/100  
 7/7 [=====] - 0s 1ms/step - loss: 133.2804 - accuracy:  
 0.6400  
 Epoch 24/100  
 7/7 [=====] - 0s 1ms/step - loss: 126.1146 - accuracy:  
 0.6300  
 Epoch 25/100

7/7 [=====] - 0s 1ms/step - loss: 227.6079 - accuracy:  
 0.6620  
 Epoch 26/100  
 7/7 [=====] - 0s 1ms/step - loss: 105.3978 - accuracy:  
 0.6460  
 Epoch 27/100  
 7/7 [=====] - 0s 1ms/step - loss: 118.8184 - accuracy:  
 0.6620  
 Epoch 28/100  
 7/7 [=====] - 0s 1ms/step - loss: 117.8343 - accuracy:  
 0.5960  
 Epoch 29/100  
 7/7 [=====] - 0s 1ms/step - loss: 189.8777 - accuracy:  
 0.6560  
 Epoch 30/100  
 7/7 [=====] - 0s 1ms/step - loss: 221.6559 - accuracy:  
 0.6520  
 Epoch 31/100  
 7/7 [=====] - 0s 1ms/step - loss: 141.7795 - accuracy:  
 0.6540  
 Epoch 32/100  
 7/7 [=====] - 0s 1ms/step - loss: 98.8320 - accuracy:  
 0.6420  
 Epoch 33/100  
 7/7 [=====] - 0s 2ms/step - loss: 98.3862 - accuracy:  
 0.6780  
 Epoch 34/100  
 7/7 [=====] - 0s 2ms/step - loss: 99.5913 - accuracy:  
 0.6820  
 Epoch 35/100  
 7/7 [=====] - 0s 2ms/step - loss: 126.5798 - accuracy:  
 0.6580  
 Epoch 36/100  
 7/7 [=====] - 0s 1ms/step - loss: 167.7500 - accuracy:  
 0.6860  
 Epoch 37/100  
 7/7 [=====] - 0s 1ms/step - loss: 288.4189 - accuracy:  
 0.6340  
 Epoch 38/100  
 7/7 [=====] - 0s 2ms/step - loss: 646.2736 - accuracy:  
 0.6900  
 Epoch 39/100  
 7/7 [=====] - 0s 2ms/step - loss: 125.7064 - accuracy:  
 0.6120  
 Epoch 40/100  
 7/7 [=====] - 0s 1ms/step - loss: 174.7963 - accuracy:  
 0.6540  
 Epoch 41/100

```

7/7 [=====] - 0s 1ms/step - loss: 138.9755 - accuracy:
0.6400
Epoch 42/100
7/7 [=====] - 0s 1ms/step - loss: 193.0526 - accuracy:
0.7120
Epoch 43/100
7/7 [=====] - 0s 1ms/step - loss: 207.9409 - accuracy:
0.6700
Epoch 44/100
7/7 [=====] - 0s 1ms/step - loss: 63.1655 - accuracy:
0.6520
Epoch 45/100
7/7 [=====] - 0s 1ms/step - loss: 91.2643 - accuracy:
0.6800
Epoch 46/100
7/7 [=====] - 0s 1ms/step - loss: 55.4114 - accuracy:
0.6900
Epoch 47/100
7/7 [=====] - 0s 1ms/step - loss: 130.1207 - accuracy:
0.6960
Epoch 48/100
7/7 [=====] - 0s 1ms/step - loss: 81.7118 - accuracy:
0.7040
Epoch 49/100
7/7 [=====] - 0s 1ms/step - loss: 123.6458 - accuracy:
0.6940
Epoch 50/100
7/7 [=====] - 0s 1ms/step - loss: 133.5079 - accuracy:
0.6900
Epoch 51/100
7/7 [=====] - 0s 1ms/step - loss: 92.7610 - accuracy:
0.6940
Epoch 52/100
7/7 [=====] - 0s 1ms/step - loss: 87.0867 - accuracy:
0.6940
Epoch 53/100
7/7 [=====] - 0s 2ms/step - loss: 148.1695 - accuracy:
0.7180
Epoch 54/100
7/7 [=====] - 0s 2ms/step - loss: 85.7722 - accuracy:
0.6920
Epoch 55/100
7/7 [=====] - 0s 2ms/step - loss: 54.7273 - accuracy:
0.7140
Epoch 56/100
7/7 [=====] - 0s 1000us/step - loss: 72.1869 -
accuracy: 0.6840
Epoch 57/100

```



```

7/7 [=====] - 0s 1ms/step - loss: 97.4154 - accuracy:
0.7120
Epoch 58/100
7/7 [=====] - 0s 1000us/step - loss: 81.3191 -
accuracy: 0.7040
Epoch 59/100
7/7 [=====] - 0s 1ms/step - loss: 91.3966 - accuracy:
0.7120
Epoch 60/100
7/7 [=====] - 0s 1ms/step - loss: 40.3106 - accuracy:
0.7080
Epoch 61/100
7/7 [=====] - 0s 1ms/step - loss: 39.9780 - accuracy:
0.7140
Epoch 62/100
7/7 [=====] - 0s 1ms/step - loss: 57.8788 - accuracy:
0.7080
Epoch 63/100
7/7 [=====] - 0s 1ms/step - loss: 164.5726 - accuracy:
0.7260
Epoch 64/100
7/7 [=====] - 0s 1ms/step - loss: 126.4698 - accuracy:
0.6920
Epoch 65/100
7/7 [=====] - 0s 1ms/step - loss: 177.5610 - accuracy:
0.7280
Epoch 66/100
7/7 [=====] - 0s 1ms/step - loss: 205.7392 - accuracy:
0.6940
Epoch 67/100
7/7 [=====] - 0s 1ms/step - loss: 60.7550 - accuracy:
0.6960
Epoch 68/100
7/7 [=====] - 0s 1ms/step - loss: 75.2348 - accuracy:
0.6960
Epoch 69/100
7/7 [=====] - 0s 1ms/step - loss: 76.4552 - accuracy:
0.6960
Epoch 70/100
7/7 [=====] - 0s 2ms/step - loss: 128.4090 - accuracy:
0.7180
Epoch 71/100
7/7 [=====] - 0s 1ms/step - loss: 92.6421 - accuracy:
0.7120
Epoch 72/100
7/7 [=====] - 0s 1000us/step - loss: 106.2934 -
accuracy: 0.7000
Epoch 73/100

```

7/7 [=====] - 0s 1ms/step - loss: 80.4447 - accuracy:  
0.7140  
Epoch 74/100  
7/7 [=====] - 0s 2ms/step - loss: 73.0985 - accuracy:  
0.6860  
Epoch 75/100  
7/7 [=====] - 0s 1ms/step - loss: 80.0589 - accuracy:  
0.7140  
Epoch 76/100  
7/7 [=====] - 0s 1ms/step - loss: 99.7002 - accuracy:  
0.7120  
Epoch 77/100  
7/7 [=====] - 0s 1ms/step - loss: 116.1769 - accuracy:  
0.7380  
Epoch 78/100  
7/7 [=====] - 0s 1ms/step - loss: 155.7042 - accuracy:  
0.6820  
Epoch 79/100  
7/7 [=====] - 0s 1ms/step - loss: 139.0850 - accuracy:  
0.6960  
Epoch 80/100  
7/7 [=====] - 0s 2ms/step - loss: 429.3281 - accuracy:  
0.7400  
Epoch 81/100  
7/7 [=====] - 0s 2ms/step - loss: 315.3177 - accuracy:  
0.6220  
Epoch 82/100  
7/7 [=====] - 0s 2ms/step - loss: 293.5067 - accuracy:  
0.7460  
Epoch 83/100  
7/7 [=====] - 0s 1ms/step - loss: 174.0258 - accuracy:  
0.6800  
Epoch 84/100  
7/7 [=====] - 0s 1ms/step - loss: 100.3365 - accuracy:  
0.6720  
Epoch 85/100  
7/7 [=====] - 0s 1ms/step - loss: 148.6979 - accuracy:  
0.7000  
Epoch 86/100  
7/7 [=====] - 0s 1ms/step - loss: 32.4058 - accuracy:  
0.6780  
Epoch 87/100  
7/7 [=====] - 0s 1ms/step - loss: 66.3409 - accuracy:  
0.7200  
Epoch 88/100  
7/7 [=====] - 0s 1ms/step - loss: 104.2426 - accuracy:  
0.7280  
Epoch 89/100

```

7/7 [=====] - 0s 1ms/step - loss: 160.2844 - accuracy:
0.6800
Epoch 90/100
7/7 [=====] - 0s 1ms/step - loss: 148.5122 - accuracy:
0.6760
Epoch 91/100
7/7 [=====] - 0s 1ms/step - loss: 209.4482 - accuracy:
0.7480
Epoch 92/100
7/7 [=====] - 0s 1ms/step - loss: 280.5800 - accuracy:
0.6860
Epoch 93/100
7/7 [=====] - 0s 1ms/step - loss: 215.6759 - accuracy:
0.7180
Epoch 94/100
7/7 [=====] - 0s 1ms/step - loss: 374.7006 - accuracy:
0.7380
Epoch 95/100
7/7 [=====] - 0s 1ms/step - loss: 193.8938 - accuracy:
0.6100
Epoch 96/100
7/7 [=====] - 0s 1ms/step - loss: 190.0665 - accuracy:
0.7280
Epoch 97/100
7/7 [=====] - 0s 1ms/step - loss: 37.0257 - accuracy:
0.7080
Epoch 98/100
7/7 [=====] - 0s 917us/step - loss: 131.2587 -
accuracy: 0.7180
Epoch 99/100
7/7 [=====] - 0s 1ms/step - loss: 245.2310 - accuracy:
0.6780
Epoch 100/100
7/7 [=====] - 0s 1ms/step - loss: 185.2444 - accuracy:
0.6940
Best: 0.754058 using {'activation': 'relu', 'batch_size': 80, 'epochs': 100,
'learning_rate': 0.001}

```

After model optimization we are able to improve accuracy from 73% to 75%, which is still an improvement. However, the model is still not very accurate, a good reason might be the imbalance of the data, which is something to keep in mind, but it's also the nature of fraud.